Data Manipulation | Exercises | Pandas | Python

101 Pandas Exercises for Data Analysis

April 27, 2018

by Selva Prabhakaran

Feedback

101 python pandas exercises are designed to challenge your logical muscle and to help internalize da manipulation with python's favorite package for data analysis. The questions are of 3 levels of difficulties with L1 being the easiest to L3 being the hardest.

Get FREE pass to my next webinar where I teach how to approach a real 'Netflix' business problem, and how to transition to a successful data science career.

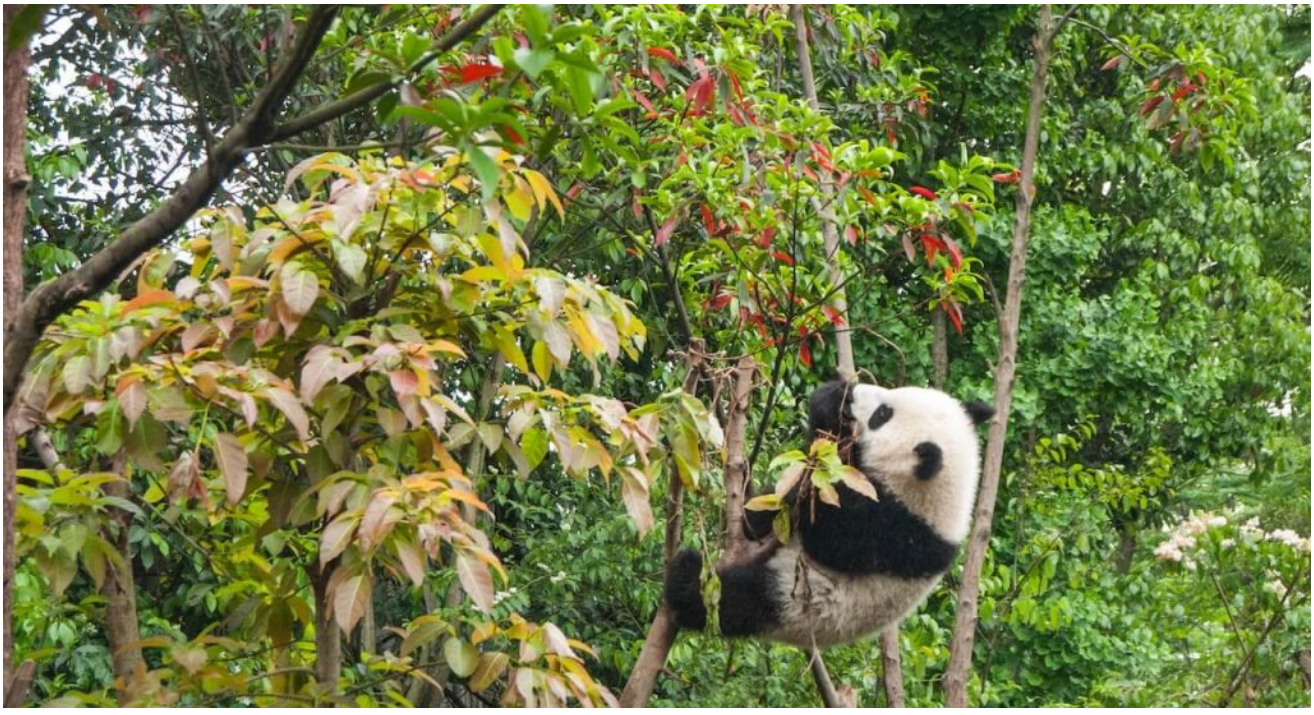

101 Pandas Exercises. Photo by Chester Ho. You might also like to practice the 101 NumPy exercises, they are often used together.

1. How to import pandas and check the version?

Show Solution ›

2. How to create a series from a list, numpy array and dict?

Create a pandas series from each of the items below: a list, numpy and a dictionary Input

```
import numpy as np
mylist = list('abcedfghijklmnopqrstuvwxyz')
myarr = np.arange(26)
mydict = dict(zip(mylist, myarr))
```

Create Powerful Visualizations in Python: Sign up for my free 9-day-email course here.

Show Solution ›

3. How to convert the index of a series into a column of a dataframe?

Difficulty Level: L1 Convert the series `ser` into a dataframe with its index as another column on the dataframe. Input

```
mylist = list('abcedfghijklmnopqrstuvwxyz')
myarr = np.arange(26)
mydict = dict(zip(mylist, myarr))
ser = pd.Series(mydict)
```

Show Solution ›

4. How to combine many series to form a dataframe?

Difficulty Level: L1 Combine ser1 and ser2 to form a dataframe. Input

```
import numpy as np
ser1 = pd.Series(list('abcedfghijklmnopqrstuvwxyz'))
ser2 = pd.Series(np.arange(26))
```

Show Solution ›

5. How to assign name to the series' index?

Difficulty Level: L1 Give a name to the series `ser` calling it 'alphabets'. Input

```
ser = pd.Series(list('abcedfghijklmnopqrstuvwxyz'))
```

Show Solution ›

6. How to get the items of series A not present in series B?

Difficulty Level: L2 From `ser1` remove items present in `ser2` .

```
ser1 = pd.Series([1, 2, 3, 4, 5])
ser2 = pd.Series([4, 5, 6, 7, 8])
```

Show Solution ›

7. How to get the items not common to both series A and series B?

Difficulty Level: L2 Get all items of `ser1` and `ser2` not common to both. Input

```
ser1 = pd.Series([1, 2, 3, 4, 5])
ser2 = pd.Series([4, 5, 6, 7, 8])
```

Show Solution ›

8. How to get the minimum, 25th percentile, median, 75th, and max of a numeric series?

Difficuty Level: L2 Compute the minimum, 25th percentile, median, 75th, and maximum of `ser` . Input

```
ser = pd.Series(np.random.normal(10, 5, 25))
```

Show Solution ›

9. How to get frequency counts of unique items of a series?

Difficulty Level: L1 Calculte the frequency counts of each unique value `ser`. Input

```
ser = pd.Series(np.take(list('abcdefgh'), np.random.randint(8, size=30)))
```

Show Solution ›

10. How to keep only top 2 most frequent values as it is and replace everything else as 'Other'?

Difficulty Level: L2 From `ser`, keep the top 2 most frequent items as it is and replace everything else as 'Other'. Input

```
np.random.RandomState(100)
ser = pd.Series(np.random.randint(1, 5, [12]))
```

Show Solution ›

11. How to bin a numeric series to 10 groups of equal size?

Difficulty Level: L2 Bin the series `ser` into 10 equal deciles and replace the values with the bin name. Input

```
ser = pd.Series(np.random.random(20))
```

Desired Output

```
 # First 5 items
0    7th
1    9th
2    7th
3    3rd
4    8th
dtype: category
Categories (10, object): [1st < 2nd < 3rd < 4th ... 7th < 8th < 9th < 10th]
```

Show Solution ›

12. How to convert a numpy array to a dataframe of given shape? (L1)

Difficulty Level: L1 Reshape the series `ser` into a dataframe with 7 rows and 5 columns Input

```
ser = pd.Series(np.random.randint(1, 10, 35))
```

Show Solution ›

13. How to find the positions of numbers that are multiples of 3 from a series?

Difficulty Level: L2 Find the positions of numbers that are multiples of 3 from `ser` . Input

```
ser = pd.Series(np.random.randint(1, 10, 7))
```

Show Solution ›

14. How to extract items at given positions from a series

Difficulty Level: L1 From `ser` , extract the items at positions in list `pos` . Input

```
ser = pd.Series(list('abcdefghijklmnopqrstuvwxyz'))
pos = [0, 4, 8, 14, 20]
```

Show Solution ›

15. How to stack two series vertically and horizontally ?

Difficulty Level: L1 Stack ser1 and ser2 vertically and horizontally (to form a dataframe). Input

```
ser1 = pd.Series(range(5))
ser2 = pd.Series(list('abcde'))
```

Show Solution ›

16. How to get the positions of items of series A in another series B?

Difficulty Level: L2 Get the positions of items of `ser2` in `ser1` as a list. Input

```
ser1 = pd.Series([10, 9, 6, 5, 3, 1, 12, 8, 13])
ser2 = pd.Series([1, 3, 10, 13])
```

Show Solution ›

17. How to compute the mean squared error on a truth and predicted series?

Difficulty Level: L2 Compute the mean squared error of `truth` and `pred` series. Input

```
truth = pd.Series(range(10))
pred = pd.Series(range(10)) + np.random.random(10)
```

Show Solution ›

18. How to convert the first character of each element in a series to uppercase?

Difficulty Level: L2 Change the first character of each word to upper case in each word of `ser`.

```
ser = pd.Series(['how', 'to', 'kick', 'ass?'])
```

Show Solution ›

19. How to calculate the number of characters in each word in a series?

Difficulty Level: L2 Input

```
ser = pd.Series(['how', 'to', 'kick', 'ass?'])
```

Show Solution ›

20. How to compute difference of differences between consequtive numbers of a series?

Difficulty Level: L1 Difference of differences between the consequtive numbers of `ser` . Input

```
ser = pd.Series([1, 3, 6, 10, 15, 21, 27, 35])
```

Desired Output

```
[nan, 2.0, 3.0, 4.0, 5.0, 6.0, 6.0, 8.0]
[nan, nan, 1.0, 1.0, 1.0, 1.0, 0.0, 2.0]
```

Show Solution ›

21. How to convert a series of date-strings to a timeseries?

Difficiulty Level: L2 Input

```
ser = pd.Series(['01 Jan 2010', '02-02-2011', '20120303', '2013/04/04', '2014-05-05', '2015-06-06T12:20
```

Desired Output

```
0   2010-01-01 00:00:00
1   2011-02-02 00:00:00
2   2012-03-03 00:00:00
3   2013-04-04 00:00:00
4   2014-05-05 00:00:00
5   2015-06-06 12:20:00
dtype: datetime64[ns]
```

Show Solution ›

22. How to get the day of month, week number, day of year and day of week from a series of date strings?

Difficiulty Level: L2 Get the day of month, week number, day of year and day of week from `ser` . Input

```
ser = pd.Series(['01 Jan 2010', '02-02-2011', '20120303', '2013/04/04', '2014-05-05', '2015-06-06T12:20
```

Desired output

```
 Date:  [1, 2, 3, 4, 5, 6]
Week number:  [53, 5, 9, 14, 19, 23]
Day num of year:  [1, 33, 63, 94, 125, 157]
Day of week:  ['Friday', 'Wednesday', 'Saturday', 'Thursday', 'Monday', 'Saturday']
```

Show Solution ›

23. How to convert year-month string to dates corresponding to the 4th day of the month?

Difficiulty Level: L2 Change `ser` to dates that start with 4th of the respective months. Input

```
ser = pd.Series(['Jan 2010', 'Feb 2011', 'Mar 2012'])
```

Desired Output

```
0   2010-01-04
1   2011-02-04
2   2012-03-04
dtype: datetime64[ns]
```

Show Solution ›

24. How to filter words that contain atleast 2 vowels from a series?

Difficiulty Level: L3 From `ser` , extract words that contain atleast 2 vowels. Input

```
ser = pd.Series(['Apple', 'Orange', 'Plan', 'Python', 'Money'])
```

Desired Output

```
 0     Apple
1    Orange
4     Money
dtype: object
```

Show Solution ›

25. How to filter valid emails from a series?

Difficiulty Level: L3 Extract the valid emails from the series `emails`. The regex pattern for valid emails is provided as reference. Input

```
 emails = pd.Series(['buying books at amazom.com', 'rameses@egypt.com', 'matt@t.co', 'narendra@modi.com'
pattern ='[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\\.[A-Za-z]{2,4}'
```

Desired Output

```
 1    rameses@egypt.com
2            matt@t.co
3    narendra@modi.com
dtype: object
```

Show Solution ›

26. How to get the mean of a series grouped by another series?

Difficiulty Level: L2 Compute the mean of `weights` of each `fruit`. Input

```
 fruit = pd.Series(np.random.choice(['apple', 'banana', 'carrot'], 10))
weights = pd.Series(np.linspace(1, 10, 10))
print(weight.tolist())
print(fruit.tolist())
#> [1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0]
#> ['banana', 'carrot', 'apple', 'carrot', 'carrot', 'apple', 'banana', 'carrot', 'apple', 'carrot']
```

Desired output

```
 # values can change due to randomness
apple     6.0
banana    4.0
```

```
carrot    5.8
dtype: float64
```

Show Solution ›

27. How to compute the euclidean distance between two series?

Difficiulty Level: L2 Compute the euclidean distance between series (points) p and q, without using a packaged formula. Input

```
p = pd.Series([1, 2, 3, 4, 5, 6, 7, 8, 9, 10])
q = pd.Series([10, 9, 8, 7, 6, 5, 4, 3, 2, 1])
```

Desired Output

```
18.165
```

Show Solution ›

28. How to find all the local maxima (or peaks) in a numeric series?

Difficiulty Level: L3 Get the positions of peaks (values surrounded by smaller values on both sides) in `ser`. Input

```
ser = pd.Series([2, 10, 3, 4, 9, 10, 2, 7, 3])
```

Desired output

```
array([1, 5, 7])
```

Show Solution ›

29. How to replace missing spaces in a string with the least frequent character?

Replace the spaces in `my_str` with the least frequent character. Difficiulty Level: L2 Input

```
my_str = 'dbc deb abed gade'
```

Desired Output

```
'dbccdebcabedcgade'  # least frequent is 'c'
```

Show Solution ›

30. How to create a TimeSeries starting '2000-01-01' and 10 weekends (saturdays) after that having random numbers as values?

Difficiulty Level: L2 Desired output

```
 # values can be random
2000-01-01    4
2000-01-08    1
2000-01-15    8
2000-01-22    4
2000-01-29    4
2000-02-05    2
2000-02-12    4
2000-02-19    9
2000-02-26    6
2000-03-04    6
```

Show Solution ›

31. How to fill an intermittent time series so all missing dates show up with values of previous non-missing date?

Difficiulty Level: L2 `ser` has missing dates and values. Make all missing dates appear and fill up with value from previous date. Input

```
 ser = pd.Series([1,10,3,np.nan], index=pd.to_datetime(['2000-01-01', '2000-01-03', '2000-01-06', '2000-
print(ser)
#> 2000-01-01     1.0
#> 2000-01-03    10.0
#> 2000-01-06     3.0
#> 2000-01-08     NaN
```

```
#> dtype: float64
```

Desired Output

```
 2000-01-01     1.0
2000-01-02     1.0
2000-01-03    10.0
2000-01-04    10.0
2000-01-05    10.0
2000-01-06     3.0
2000-01-07     3.0
2000-01-08     NaN
```

Show Solution ›

32. How to compute the autocorrelations of a numeric series?

Difficiulty Level: L3 Compute autocorrelations for the first 10 lags of `ser`. Find out which lag has the largest correlation. Input

```
ser = pd.Series(np.arange(20) + np.random.normal(1, 10, 20))
```

Desired output

```
 # values will change due to randomness
[0.29999999999999999, -0.11, -0.17000000000000001, 0.46000000000000002, 0.28000000000000003, -0.040000000
Lag having highest correlation:  9
```

Show Solution ›

33. How to import only every nth row from a csv file to create a dataframe?

Difficiulty Level: L2 Import every 50th row of BostonHousing dataset as a dataframe. Show Solution ›

34. How to change column values when importing csv to a dataframe?

Difficulty Level: L2 Import the boston housing dataset, but while importing change the `'medv'` (median house value) column so that values < 25 becomes 'Low' and > 25 becomes 'High'. Show Solution ›

35. How to create a dataframe with rows as strides from a

given series?

Difficiulty Level: L3 Input

```
L = pd.Series(range(15))
```

Desired Output

```
array([[ 0,  1,  2,  3],
       [ 2,  3,  4,  5],
       [ 4,  5,  6,  7],
       [ 6,  7,  8,  9],
       [ 8,  9, 10, 11],
       [10, 11, 12, 13]])
```

Show Solution ›

36. How to import only specified columns from a csv file?

Difficulty Level: L1 Import 'crim' and 'medv' columns of the BostonHousing dataset as a dataframe.

Show Solution ›

37. How to get the *n*rows, *n*columns, datatype, summary stats of each column of a dataframe? Also get the array and list equivalent.

Difficulty Level: L2 Get the number of rows, columns, datatype and summary statistics of each column of the Cars93 dataset. Also get the numpy array and list equivalent of the dataframe. Show Solution ›

38. How to extract the row and column number of a particular cell with given criterion?

Difficulty Level: L1 Input

```
df = pd.read_csv('https://raw.githubusercontent.com/selva86/datasets/master/Cars93_miss.csv')
```

Which manufacturer, model and type has the highest `Price` ? What is the row and column number of the cell with the highest `Price` value? Show Solution ›

39. How to rename a specific columns in a dataframe?

Difficulty Level: L2 Rename the column `Type` as `CarType` in `df` and replace the '.' in column names with '_'. Input

```
df = pd.read_csv('https://raw.githubusercontent.com/selva86/datasets/master/Cars93_miss.csv')
print(df.columns)
#> Index(['Manufacturer', 'Model', 'Type', 'Min.Price', 'Price', 'Max.Price',
#>        'MPG.city', 'MPG.highway', 'AirBags', 'DriveTrain', 'Cylinders',
#>        'EngineSize', 'Horsepower', 'RPM', 'Rev.per.mile', 'Man.trans.avail',
#>        'Fuel.tank.capacity', 'Passengers', 'Length', 'Wheelbase', 'Width',
#>        'Turn.circle', 'Rear.seat.room', 'Luggage.room', 'Weight', 'Origin',
#>        'Make'],
#>       dtype='object')
```

Desired Solution

```
print(df.columns)
#> Index(['Manufacturer', 'Model', 'CarType', 'Min_Price', 'Price', 'Max_Price',
#>        'MPG_city', 'MPG_highway', 'AirBags', 'DriveTrain', 'Cylinders',
#>        'EngineSize', 'Horsepower', 'RPM', 'Rev_per_mile', 'Man_trans_avail',
#>        'Fuel_tank_capacity', 'Passengers', 'Length', 'Wheelbase', 'Width',
#>        'Turn_circle', 'Rear_seat_room', 'Luggage_room', 'Weight', 'Origin',
#>        'Make'],
#>       dtype='object')
```

Show Solution ›

40. How to check if a dataframe has any missing values?

Difficulty Level: L1 Check if `df` has any missing values. Input

```
df = pd.read_csv('https://raw.githubusercontent.com/selva86/datasets/master/Cars93_miss.csv')
```

Show Solution ›

41. How to count the number of missing values in each column?

Difficulty Level: L2 Count the number of missing values in each column of `df`. Which column has the maximum number of missing values? Input

```
df = pd.read_csv('https://raw.githubusercontent.com/selva86/datasets/master/Cars93_miss.csv')
```

Show Solution ›

42. How to replace missing values of multiple numeric columns with the mean?

Difficulty Level: L2 Replace missing values in `Min.Price` and `Max.Price` columns with their respective mean. Input

```
df = pd.read_csv('https://raw.githubusercontent.com/selva86/datasets/master/Cars93_miss.csv')
```

Show Solution ›

43. How to use apply function on existing columns with global variables as additional arguments?

Difficulty Level: L3 In `df`, use `apply` method to replace the missing values in `Min.Price` with the column's mean and those in `Max.Price` with the column's median. Input

```
df = pd.read_csv('https://raw.githubusercontent.com/selva86/datasets/master/Cars93_miss.csv')
```

Use Hint from StackOverflow Show Solution ›

44. How to select a specific column from a dataframe as a dataframe instead of a series?

Difficulty Level: L2 Get the first column (`a`) in `df` as a dataframe (rather than as a Series). Input

```
df = pd.DataFrame(np.arange(20).reshape(-1, 5), columns=list('abcde'))
```

Show Solution ›

45. How to change the order of columns of a dataframe?

Difficulty Level: L3 Actually 3 questions.

1. In `df`, interchange columns `'a'` and `'c'`.
2. Create a generic function to interchange two columns, without hardcoding column names.
3. Sort the columns in reverse alphabetical order, that is colume `'e'` first through column `'a'` last.

Input

```
df = pd.DataFrame(np.arange(20).reshape(-1, 5), columns=list('abcde'))
```

Show Solution ›

46. How to set the number of rows and columns displayed in the output?

Difficulty Level: L2 Change the pamdas display settings on printing the dataframe `df` it shows a maximum of 10 rows and 10 columns. Input

```
df = pd.read_csv('https://raw.githubusercontent.com/selva86/datasets/master/Cars93_miss.csv')
```

Show Solution ›

47. How to format or suppress scientific notations in a pandas dataframe?

Difficulty Level: L2 Suppress scientific notations like 'e-03' in `df` and print upto 4 numbers after decimal. Input

```
 df = pd.DataFrame(np.random.random(4)**10, columns=['random'])
df
#>          random
#> 0  3.474280e-03
#> 1  3.951517e-05
#> 2  7.469702e-02
#> 3  5.541282e-28
```

Desired Output

```
 #>    random
#> 0  0.0035
#> 1  0.0000
#> 2  0.0747
```

```
#> 3  0.0000
```

Show Solution ›

48. How to format all the values in a dataframe as percentages?

Difficulty Level: L2 Format the values in column `'random'` of `df` as percentages. Input

```
df = pd.DataFrame(np.random.random(4), columns=['random'])
df
#>       random
#> 0    .689723
#> 1    .957224
#> 2    .159157
#> 3    .21082
```

Desired Output

```
#>       random
#> 0    68.97%
#> 1    95.72%
#> 2    15.91%
#> 3    2.10%
```

Show Solution ›

49. How to filter every nth row in a dataframe?

Difficulty Level: L1 From `df`, filter the `'Manufacturer'`, `'Model'` and `'Type'` for every 20th row starting from 1st (row 0). Input

```
df = pd.read_csv('https://raw.githubusercontent.com/selva86/datasets/master/Cars93_miss.csv')
```

Show Solution ›

50. How to create a primary key index by combining relevant columns?

Difficulty Level: L2 In `df`, Replace `NaN`s with 'missing' in columns `'Manufacturer'`, `'Model'` and `'Type'` and create a index as a combination of these three columns and check if the index is a primary key. Input

```
df = pd.read_csv('https://raw.githubusercontent.com/selva86/datasets/master/Cars93_miss.csv', usecols=[
```

Desired Output

	Manufacturer	Model	Type	Min.Price	Max.Price
Acura_Integra_Small	Acura	Integra	Small	12.9	18.8
missing_Legend_Midsize	missing	Legend	Midsize	29.2	38.7
Audi_90_Compact	Audi	90	Compact	25.9	32.3
Audi_100_Midsize	Audi	100	Midsize	NaN	44.6
BMW_535i_Midsize	BMW	535i	Midsize	NaN	NaN

Show Solution ›

51. How to get the row number of the nth largest value in a column?

Difficulty Level: L2 Find the row position of the 5th largest value of column `'a'` in `df`. Input

```
df = pd.DataFrame(np.random.randint(1, 30, 30).reshape(10,-1), columns=list('abc'))
```

Show Solution ›

52. How to find the position of the nth largest value greater than a given value?

Difficulty Level: L2 In `ser`, find the position of the 2nd largest value greater than the mean. Input

```
ser = pd.Series(np.random.randint(1, 100, 15))
```

Show Solution ›

53. How to get the last n rows of a dataframe with row sum > 100?

Difficulty Level: L2 Get the last two rows of `df` whose row sum is greater than 100.

```
df = pd.DataFrame(np.random.randint(10, 40, 60).reshape(-1, 4))
```

Show Solution

54. How to find and cap outliers from a series or dataframe column?

Difficulty Level: L2 Replace all values of `ser` in the lower 5%ile and greater than 95%ile with respective 5th and 95th %ile value. Input

```
ser = pd.Series(np.logspace(-2, 2, 30))
```

Show Solution

55. How to reshape a dataframe to the largest possible square after removing the negative values?

Difficulty Level: L3 Reshape `df` to the largest possible square with negative values removed. Drop the smallest values if need be. The order of the positive numbers in the result should remain the same as the original. Input

```
df = pd.DataFrame(np.random.randint(-20, 50, 100).reshape(10,-1))
```

Show Solution

56. How to swap two rows of a dataframe?

Difficulty Level: L2 Swap rows 1 and 2 in `df`. Input

```
df = pd.DataFrame(np.arange(25).reshape(5, -1))
```

Show Solution

57. How to reverse the rows of a dataframe?

Difficulty Level: L2 Reverse all the rows of dataframe `df`. Input

```
df = pd.DataFrame(np.arange(25).reshape(5, -1))
```

Show Solution

58. How to create one-hot encodings of a categorical variable (dummy variables)?

Difficulty Level: L2 Get one-hot encodings for column `'a'` in the dataframe `df` and append it as columns. Input

```
df = pd.DataFrame(np.arange(25).reshape(5,-1), columns=list('abcde'))
```

```
   a   b   c   d   e
0  0   1   2   3   4
1  5   6   7   8   9
2  10  11  12  13  14
3  15  16  17  18  19
4  20  21  22  23  24
```

Output

```
    0  5  10  15  20   b   c   d   e
0   1  0   0   0   0   1   2   3   4
1   0  1   0   0   0   6   7   8   9
2   0  0   1   0   0  11  12  13  14
3   0  0   0   1   0  16  17  18  19
4   0  0   0   0   1  21  22  23  24
```

Show Solution ›

59. Which column contains the highest number of row-wise maximum values?

Difficulty Level: L2 Obtain the column name with the highest number of row-wise maximum's in `df`.

```
df = pd.DataFrame(np.random.randint(1,100, 40).reshape(10, -1))
```

Show Solution ›

60. How to create a new column that contains the row number of nearest column by euclidean distance?

Create a new column such that, each row contains the row number of nearest row-record by euclidean distance. Difficulty Level: L3 Input

```
df = pd.DataFrame(np.random.randint(1,100, 40).reshape(10, -1), columns=list('pqrs'), index=list('abcde
df
#     p   q   r   s
# a  57  77  13  62
# b  68   5  92  24
# c  74  40  18  37
# d  80  17  39  60
# e  93  48  85  33
# f  69  55   8  11
# g  39  23  88  53
# h  63  28  25  61
# i  18   4  73   7
# j  79  12  45  34
```

Desired Output

```
df
#    p   q   r   s nearest_row   dist
# a  57  77  13  62           i  116.0
# b  68   5  92  24           a  114.0
# c  74  40  18  37           i   91.0
# d  80  17  39  60           i   89.0
# e  93  48  85  33           i   92.0
# f  69  55   8  11           g  100.0
# g  39  23  88  53           f  100.0
# h  63  28  25  61           i   88.0
# i  18   4  73   7           a  116.0
# j  79  12  45  34           a   81.0
```

Show Solution ›

61. How to know the maximum possible correlation value of each column against other columns?

Difficulty Level: L2 Compute maximum possible absolute correlation value of each column against other columns in `df` . Input

```
df = pd.DataFrame(np.random.randint(1,100, 80).reshape(8, -1), columns=list('pqrstuvwxy'), index=list('
```

Show Solution ›

62. How to create a column containing the minimum by maximum of each row?

Difficulty Level: L2 Compute the minimum-by-maximum for every row of `df`.

```
df = pd.DataFrame(np.random.randint(1,100, 80).reshape(8, -1))
```

Show Solution

63. How to create a column that contains the penultimate value in each row?

Difficulty Level: L2 Create a new column `'penultimate'` which has the second largest value of each row of `df`. Input

```
df = pd.DataFrame(np.random.randint(1,100, 80).reshape(8, -1))
```

Show Solution

64. How to normalize all columns in a dataframe?

Difficulty Level: L2

1. Normalize all columns of `df` by subtracting the column mean and divide by standard deviation.
2. Range all columns of `df` such that the minimum value in each column is 0 and max is 1.

Don't use external packages like sklearn. Input

```
df = pd.DataFrame(np.random.randint(1,100, 80).reshape(8, -1))
```

Show Solution

65. How to compute the correlation of each row with the suceeding row?

Difficulty Level: L2 Compute the correlation of each row of `df` with its succeeding row. Input

```
df = pd.DataFrame(np.random.randint(1,100, 80).reshape(8, -1))
```

Show Solution ›

66. How to replace both the diagonals of dataframe with 0?

Difficulty Level: L2 Replace both values in both diagonals of `df` with 0. Input

```
 df = pd.DataFrame(np.random.randint(1,100, 100).reshape(10, -1))
df
#     0   1   2   3   4   5   6   7   8   9
# 0  11  46  26  44  11  62  18  70  68  26
# 1  87  71  52  50  81  43  83  39   3  59
# 2  47  76  93  77  73   2   2  16  14  26
# 3  64  18  74  22  16  37  60   8  66  39
# 4  10  18  39  98  25   8  32   6   3  29
# 5  29  91  27  86  23  84  28  31  97  10
# 6  37  71  70  65   4  72  82  89  12  97
# 7  65  22  97  75  17  10  43  78  12  77
# 8  47  57  96  55  17  83  61  85  26  86
# 9  76  80  28  45  77  12  67  80   7  63
```

Desired output

```
 #     0   1   2   3   4   5   6   7   8   9
# 0   0  46  26  44  11  62  18  70  68   0
# 1  87   0  52  50  81  43  83  39   0  59
# 2  47  76   0  77  73   2   2   0  14  26
# 3  64  18  74   0  16  37   0   8  66  39
# 4  10  18  39  98   0   0  32   6   3  29
# 5  29  91  27  86   0   0  28  31  97  10
# 6  37  71  70   0   4  72   0  89  12  97
# 7  65  22   0  75  17  10  43   0  12  77
# 8  47   0  96  55  17  83  61  85   0  86
# 9   0  80  28  45  77  12  67  80   7   0
```

Show Solution ›

67. How to get the particular group of a groupby dataframe by key?

Difficulty Level: L2 This is a question related to understanding of grouped dataframe. From `df_grouped`, get the group belonging to `'apple'` as a dataframe. Input

```
df = pd.DataFrame({'col1': ['apple', 'banana', 'orange'] * 3,
                   'col2': np.random.rand(9),
                   'col3': np.random.randint(0, 15, 9)})
```

```
df_grouped = df.groupby(['col1'])
```

```
 # Input
df = pd.DataFrame({'col1': ['apple', 'banana', 'orange'] * 3,
                   'col2': np.random.rand(9),
                   'col3': np.random.randint(0, 15, 9)})

df_grouped = df.groupby(['col1'])

# Solution 1
df_grouped.get_group('apple')

# Solution 2
for i, dff in df_grouped:
    if i == 'apple':
        print(dff)
```

```
     col1      col2  col3
0   apple  0.673434     7
3   apple  0.182348    14
6   apple  0.050457     3
```

[/expand]

68. How to get the n'th largest value of a column when grouped by another column?

Difficulty Level: L2 In `df` , find the second largest value of `'taste'` for `'banana'` Input

```
df = pd.DataFrame({'fruit': ['apple', 'banana', 'orange'] * 3,
                   'rating': np.random.rand(9),
                   'price': np.random.randint(0, 15, 9)})
```

Show Solution ›

69. How to compute grouped mean on pandas dataframe and keep the grouped column as another column (not index)?

Difficulty Level: L1 In `df` , Compute the mean `price` of every `fruit` , while keeping the `fruit` as another column instead of an index. Input

```
df = pd.DataFrame({'fruit': ['apple', 'banana', 'orange'] * 3,
                   'rating': np.random.rand(9),
                   'price': np.random.randint(0, 15, 9)})
```

Show Solution

70. How to join two dataframes by 2 columns so they have only the common rows?

Difficulty Level: L2 Join dataframes `df1` and `df2` by 'fruit-pazham' and 'weight-kilo'. Input

```
df1 = pd.DataFrame({'fruit': ['apple', 'banana', 'orange'] * 3,
                    'weight': ['high', 'medium', 'low'] * 3,
                    'price': np.random.randint(0, 15, 9)})

df2 = pd.DataFrame({'pazham': ['apple', 'orange', 'pine'] * 2,
                    'kilo': ['high', 'low'] * 3,
                    'price': np.random.randint(0, 15, 6)})
```

Show Solution

72. How to get the positions where values of two columns match?

Difficulty Level: L2 Show Solution

73. How to create lags and leads of a column in a dataframe?

Difficulty Level: L2 Create two new columns in `df`, one of which is a lag1 (shift column `a` down by 1 row) of column 'a' and the other is a lead1 (shift column `b` up by 1 row). Input

```
df = pd.DataFrame(np.random.randint(1, 100, 20).reshape(-1, 4), columns = list('abcd'))
```

```
    a   b   c   d
0  66  34  76  47
1  20  86  10  81
2  75  73  51  28
3   1   1   9  83
4  30  47  67   4
```

Desired Output

```
    a   b   c   d  a_lag1  b_lead1
0  66  34  76  47     NaN     86.0
1  20  86  10  81    66.0     73.0
2  75  73  51  28    20.0      1.0
3   1   1   9  83    75.0     47.0
4  30  47  67   4     1.0      NaN
```

Show Solution

74. How to get the frequency of unique values in the entire dataframe?

Difficulty Level: L2 Get the frequency of unique values in the entire dataframe `df`. Input

```
df = pd.DataFrame(np.random.randint(1, 10, 20).reshape(-1, 4), columns = list('abcd'))
```

Show Solution

75. How to split a text column into two separate columns?

Difficulty Level: L2 Split the string column in `df` to form a dataframe with 3 columns as shown. Input

```
df = pd.DataFrame(["STD, City    State",
"33, Kolkata    West Bengal",
"44, Chennai    Tamil Nadu",
"40, Hyderabad    Telengana",
"80, Bangalore    Karnataka"], columns=['row'])

print(df)
#>                         row
#> 0          STD, City\tState
#> 1  33, Kolkata\tWest Bengal
#> 2   44, Chennai\tTamil Nadu
#> 3  40, Hyderabad\tTelengana
#> 4  80, Bangalore\tKarnataka
```

Desired Output

```
  0 STD        City        State
1  33     Kolkata  West Bengal
2  44     Chennai   Tamil Nadu
3  40   Hyderabad    Telengana
4  80   Bangalore    Karnataka
```

Show Solution To be continued . .

Selva Prabhakaran

Selva is the Chief Author and Editor of Machine Learning Plus, with 4 Million+ readership. He has authored courses and books with100K+ students, and is the Principal Data Scientist of a global firm.

Previous Article
LDA in Python – How to grid search best topic models?

Next Article
Feature Selection – Ten Effective Techniq...

referred to as 'Anonymous function' is same as a ...

Analysis (PCA) is an algorithm to transform ...

to all major NLP problems from basic use of packages ...

module for working with dates in python. It ...

to v dist

Debaditya Nath • 23 days ago

please update the answers for some of them, np.argwhere dosent work on series anymore

Reply • Share ›

Debaditya Nath • 23 days ago

alternate for #34 question

```
df = pd.read_csv("https://raw.githubusercontent.com/selva86/datasets/master/BostonHousing.csv")
for ind in df.index:
if df.loc[ind, "medv"] <= 25:
df.loc[ind, "medv"] = "Low"
elif df.loc[ind, "medv"] > 25:
df.loc[ind, "medv"] = "High"
df
```

Reply • Share ›

Debaditya Nath • 25 days ago

whats up with the series words in the 18th and 19th one

Reply • Share ›

Debaditya Nath • 25 days ago

22nd one `dt.weekday_name` dosent work, you have to use `dt.day_name()` with the brackets

Reply • Share ›

Debaditya Nath • 25 days ago

many of them have have solution that require nothing but just python
for example i did this for the 24th one

```
emails = pd.Series(['buying books at amazom.com', 'rameses@egypt.com', 'matt@t.co', 'narendra@modi.com'])
import re
pattern = "[a-zA-Z_._%+-0-9]+@[a-zA-Z_0-9.-]*\.[a-zA-Z]*"
final = {}
for index, i in enumerate(emails):
match = re.findall(pattern, i)
try:
match[0]
except:
match.append("")
if i == match[0]:
final[index] = i
pd.Series(final)
```

Reply • Share ›

LanternD • a year ago • edited

#22: `weekday_name` was deprecated. It is replaced by `df.dt.day_name()`.

Reply • Share ›

Andy F • a year ago

#16

pd.Series([ser2.apply(lambda x: np.argwhere(ser1.values == x))])

Reply • Share ›

Andy F ➚ Andy F • a year ago

...or even just:

ser2.apply(lambda x: np.argwhere(ser1.values == x))

Reply • Share ›

Andy F • a year ago

I love this!

Reply • Share ›

valleyease • a year ago

Alternate solutions:

```
#18. ser.str.title()
#19. ser.str.len()
#23. pd.to_datetime(ser) + pd.DateOffset(days=3)
#24. ser.loc[ser.str.lower().map(lambda x: len(set(x).intersection(set('aeiou'))))>1]
#25. pd.Series(emails.reindex(emails).filter(regex=pattern).index)
#28. ser.loc[(ser.diff()>0)&(ser.diff(-1)>0)].index.to_list()
#29. my_str.replace(' ', Counter(my_str.replace(' ', '')).most_common()[-1][0])
#33. pd.read_csv('https://raw.githubusercontent.com/selva86/datasets/master/BostonHousing.csv', skiprows=lambda x: x%50!=0)
#39. df.rename(columns=dict(zip(df.columns, df.columns.str.replace('.', '_')))).rename(columns = {'Type':'CarType'})
#41. df.isnull().sum().idxmax()
#51. df['a'].sort_values(ascending=True).index[5]
#52. ser.loc[ser > ser.mean()].index[1]
#53. df.loc[df.sum(axis=1)>100].iloc[-2:]
#57. df.set_index(df.index[::-1]).sort_index()
#59. df.max(axis=1).argmax()
#63. df['penultimate'] = df.apply(lambda x: sorted(set(x))[-2], axis=1)
#66. np.fill_diagonal(df.to_numpy(), 0)
```

Reply • Share ›

Andrei Radu • 2 years ago

Alternate question 53 answer:

```
np.argwhere(np.array(df).sum(axis = 1) > 100)[-2:]
```

Reply • Share ›

Andrei Radu ➚ Andrei Radu • 2 years ago • edited

#64:
2.Range all columns of df such that the minimum value in each column is 0 and max is 1.

The solution is wrong as it outputs the minimum as 1, and max is 0.
A correction I found is :

```
df.apply(lambda x:(-(x.min() - x)/(x.max()-x.min())).round(2))
```

Reply • Share ›

Andrei Radu ➚ Andrei Radu • 2 years ago

Question 51 solution is slightly wrong, it should be `df['a'].argsort()[::-1][n-1]`

Reply • Share ›

Matt • 2 years ago

#38: (Correction)

```
df.get_value(row[0], 'Price')
```

should be,

```
df.loc[row[0], 'Price']
```

Reply • Share ›

Matt • 2 years ago

#22: Last line of solution should be

```
print("Day of week: ", ser_ts.dt.day_name().tolist())
```

Reply • Share ›

Itai Seri • 2 years ago • edited

Shorter way for #59:

```
df.idxmax(axis=1).value_counts().idxmax()
```

Reply • Share ›

Haksell • 2 years ago

#23 : `pd.to_datetime(ser) + np.timedelta64(3, "D")`

Reply • Share ›

Haksell • 2 years ago

#20 : "consequtive" -> "consecutive"

Reply • Share ›

Haksell • 2 years ago

#18 : `ser.str.capitalize()`

Reply • Share ›

Haksell ➚ Haksell • 2 years ago

In the same vein for #19 : `ser.str.len()`

Reply • Share ›

Haksell • 2 years ago

For question 14, `ser[pos]` is much better.

Reply • Share ›

Haksell • 2 years ago

There is no question 71

1 Reply • Share ›

msvsr • 2 years ago

Alternate solution for 16th:
print(np.where(ser1.isin(ser2)))

Reply • Share ›

msvsr • 2 years ago

13th solution:
np.argwhere(ser.values % 3 == 0)

Reply • Share ›

msvsr • 2 years ago

solution for 14th question:
np.argwhere(ser.values % 3 == 0)

Reply • Share ›

ehsan negahbani • 2 years ago • edited

Alternate to #31, if you do not know the resample() function:

```
ind = []
vals = []
for i in range(len(ser)-1):
------ind.append(ser.index[i])
------vals.append(ser.values[i])
------gap_day = (ser.index[i+1] - ser.index[i]).days-1
------for j in range(gap_day):
---------ind.append(ser.index[i]+timedelta(days=j+1))
```

```
---------vals.append(ser.values[i])
ind.append(ser.index[-1])
vals.append(ser.values[-1])
ser2 = pd.Series(vals, index = ind)
ser2
```

Reply • Share ›

Andrea D. • 2 years ago • edited

Hi, in question 52 it looks to me you are returning the 2nd element greater than the mean, but you asked for the 2nd largest element so you should try something different like this:

```
#Take elements greater than meand
arr1 = np.argwhere(ser > ser.mean())
#Transform it into a list of int to be used with iloc
arr2 = [n[0] for n in arr1]
# Let's create a working df with the values sorted (column called 'pos') and the postion as index
df = pd.DataFrame(ser.iloc[arr2], columns=['pos'])
# Now we can take the index of the second element of the df sorted by pos in descending order
row = df['pos'].sort_values(ascending= False).index.values.astype(int)[1]
#That's our value
row
```

Thank you very much for this post, it's been really useful to me!

1 Reply • Share ›

ehsan negahbani • 2 years ago

Another solution for @25:

```
import re as re
[email for email in emails if re.findall(pattern, email)]
```

Reply • Share ›

ehsan negahbani • 2 years ago

Another solution to #24 without using "collections":

```
ser = pd.Series(['Apple', 'Orange', 'Plan', 'Python', 'Money'])
vowels = pd.Series(['a', 'e', 'i', 'o', 'u'])
[x for x in ser if len(np.intersect1d(list(x.lower()), vowels))>1]
```

Reply • Share ›

Lina • 2 years ago

Can anyone explain to me why why in #60 its max instead of min? Shouldnt it be that nearest distance is the lowest one?

1 Reply • Share ›

Andrei Radu ➚ Lina • 2 years ago

I'm wondering the same thing, I think it has to be a mistake

Reply • Share ›

Laster Fahrer ➚ Andrei Radu • a year ago

I agree. It should be min.

Reply • Share ›

Richard Croft • 3 years ago

Question 71 reminds me of platform nine and three quarters...

Reply • Share ›

Richard Croft • 3 years ago

Excellent set of question, thanks v much

Reply • Share ›

Selva Prabhakaran Mod ➚ Richard Croft • 2 years ago

Welcome :)

Reply • Share ›

Surya Teja Parnampedu • 3 years ago • edited

Alternate for #24:
ser[ser.str.count(pat=r'[aeiou]', flags=re.I) >= 2]

Alternate for #25:
emails[emails.str.match(pat=r"[A-z0-9._%+-]+@[A-z0-9.-]+\.[A-z]{2,4}")]

Alternate for #42:
df.fillna({
'Min_Price': df.Min_Price.mean(),
'Max_Price': df.Max_Price.mean()
})

Reply • Share ›

Selva Prabhakaran Mod ➚ Surya Teja Parnampedu • 2 years ago

Nice

Reply • Share ›

Bhishan Poudel • 3 years ago

Qn 25

```
emails = pd.Series(['buying books at amazom.com', 'rameses@egypt.com', 'matt@t.co', 'narendra@modi.com'])

pattern = '([A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\\.[A-Za-z]{2,4})'
print(emails.str.extract(pattern, flags=re.I))

0
0 NaN
1 rameses@egypt.com
2 matt@t.co
3 narendra@modi.com
```

Reply • Share ›

Bhishan Poudel • 3 years ago

Qn 24

probably slow, but easy.

```
ser[ser.apply(lambda x: sum(map(x.lower().count,'aeiou'))) >= 2]
```

Reply • Share ›

Selva Prabhakaran Mod ➚ Bhishan Poudel • 2 years ago

Thanks

Reply • Share ›

Bhishan Poudel • 3 years ago

Qn 23

```
pd.to_datetime("04 " + ser)
```

Reply • Share ›

Bhishan Poudel • 3 years ago

Qn 10

using value_counts (EASIER)

```
%%timeit

np.random.seed(100)
ser = pd.Series(np.random.randint(1, 5, [12]))

idx = ser.value_counts().head(2).index
ser[~ser.isin(idx)] = 'Other'
ser
1.5 ms ± 18.4 µs per loop (mean ± std. dev. of 7 runs, 1000 loops each)
```

# using counter (FASTER)
from collections import Counter

```
%%timeit
np.random.seed(100)
ser = pd.Series(np.random.randint(1, 5, [12]))
```

```
top2 = Counter(ser.values).most_common(2)
idx = [i[0] for i in top2]

ser[~ser.isin(idx)] = 'Other'
ser
1.11 ms ± 10.1 µs per loop (mean ± std. dev. of 7 runs, 1000 loops each)
```

1 ▲ | ▼ • Reply • Share ›

Bhishan Poudel ➚ Bhishan Poudel • 2 years ago • edited

I revisited it long after.

▲ | ▼ • Reply • Share ›

Selva Prabhakaran Mod ➚ Bhishan Poudel • 2 years ago

Thanks for sharing

▲ | ▼ • Reply • Share ›

Bhishan Poudel • 3 years ago

Qn 9

Easiest option is of course the given solution value_counts().
However, we can also do this using numpy.

```
np.random.seed(100)
ser = pd.Series([np.random.choice(list('abcdef')) for _ in range(30)])
ser.value_counts()
c 7
a 6
e 5
d 5
f 4
b 3
dtype: int64

# using numpy
u,c = np.unique(ser.values, return_counts= True)
np.array([u,c]).T
array([['a', 6],
['b', 3],
['c', 7],
['d', 5],
['e', 5],
['f', 4]], dtype=object)
```

▲ | ▼ • Reply • Share ›

Bhishan Poudel • 3 years ago

**Qn **

```
ser.describe()
```

▲ | ▼ • Reply • Share ›

Bhishan Poudel • 3 years ago

Qn 7

```
s = pd.Series(np.setxor1d(ser1.values, ser2.values))
```

▲ | ▼ • Reply • Share ›

Bhishan Poudel • 3 years ago • edited

Qn 6

```
%%timeit
ser1[~ser1.isin(ser2)]
449 µs ± 5.01 µs per loop (mean ± std. dev. of 7 runs, 1000 loops each)

%%timeit
s = pd.Series(np.setdiff1d(ser1.values, ser2.values))
87.2 µs ± 309 ns per loop (mean ± std. dev. of 7 runs, 10000 loops each)
```

▲ | ▼ • Reply • Share ›

Bhishan Poudel • 3 years ago • edited

Qn 54

```
capped_ser = np.clip(ser, *np.percentile(ser,[5,95]) )
```

Reply • Share ›

Bhishan Poudel ➚ Bhishan Poudel • 2 years ago

Edit:
capped_arr = np.clip(ser.to_numpy(), *np.percentile(ser.to_numpy(),[5,95]))

capped_ser = pd.Series(capped_arr)
capped_ser

1 Reply • Share ›

Load more comments

Related Posts

- cProfile – How to profile your python code
- Dask Tutorial – How to handle big data in Python
- Modin – How to speedup pandas
- What does Python Global Interpreter Lock – (GIL) do?
- Python Yield – What does the yield keyword do?
- Lambda Function in Python – How and When to use?
- Investor's Portfolio Optimization with Python
- datetime in Python – Simplified Guide with Clear Examples
- Python Collections – Complete Guide
- pdb – How to use Python debugger
- Python JSON – Guide
- How to use tf.function to speed up Python code in Tensorflow
- List Comprehensions in Python – My Simplified Guide
- Mahalonobis Distance – Understanding the math with examples (python)

More Articles

Python

Numpy Reshape – How to reshape arrays and what does -1 mean?

Python

Vaex – Faster Pandas Alternate in Python

Python

Modin – How to speedup pandas by changing one line of code

May 26, 2020

Enter Email*

Join

Subscribe to Machine Learning Plus for high value data science content

Blogs

Courses

Time Series Template

Terms of Use

Privacy Policy

Contact Us

Refund Policy