

Input Data Analysis and Distribution Fitting with R

Lidia Montero

September 2016

- 1 Introduction to (Univariate) Distribution Fitting
 - 1.1 Summarize data
 - 1.2 Autocorrelation Function
- 2 Plot data
 - 2.1 Histogram: Equal length intervals
- 3 List of Candidate distributions
 - 3.1 Plot pdf's for variable location = 0, scale = 1 for several shape parameter values
 - 3.2 Method `descdist()` in `fitdistrplus` package
- 4 Fit distribution
- 5 Goodness of Fit
 - 5.1 Histogram: Variable length intervals
 - 5.2 Validate matching between empirical moments and theoretical moments
 - 5.3 Goodness of Fit: Graphical Assessment
 - 5.3.1 Known shape
 - 5.3.2 UNKnown shape
 - 5.4 Goodness of Fit: Tests
 - 5.4.1 Kolmogorov-Smirnoff Test for continuous data
 - 5.4.2 Cramer von Miess Test for discrete data
 - 5.4.3 Chi Squared Test

1 Introduction to (Univariate) Distribution Fitting

I generate a sequence of 5000 numbers distributed following a Weibull distribution with:

1. $c = \text{location} = 10$ (shift from origin),
2. $b = \text{scale} = 2$ and
3. $a = \text{shape} = 1$

```
sample<- rweibull(5000, shape=1, scale = 2) + 10
```

The Weibull distribution with shape parameter a and scale parameter b has density given by

$$f(x) = (a/b) (x/b)^{a-1} \exp(- (x/b)^a) \text{ for } x > 0$$

The cumulative distribution function is $F(x) = 1 - \exp(- (x/b)^a)$ on $x > 0$.

Theoretical moments for Weibull distributions are:

1. $\text{Mean}(X) = E(X) = c + b \Gamma(1 + 1/a)$
2. $\text{Var}(X) = b^2 * (\Gamma(1 + 2/a) - (\Gamma(1 + 1/a))^2)$

1.1 Summarize data

Don't forget to validate uncorrelated sample data :

- `acf()` Autocorrelation function is fast and easy in R.
- Use `durbinWatsonTest()` for an inferential option.
- Calculate central and plain moments (up to order 4) using method `all.moments()` in `library(moments)`

```
summary(sample)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    10.00   10.59   11.41   12.03   12.83   25.92
```

```
library(moments)
all.moments(sample,4,central=F)
```

```
## [1]      1.000000     12.03262    148.89875   1906.48619  25430.78743
```

```
all.moments(sample,4,central=T)
```

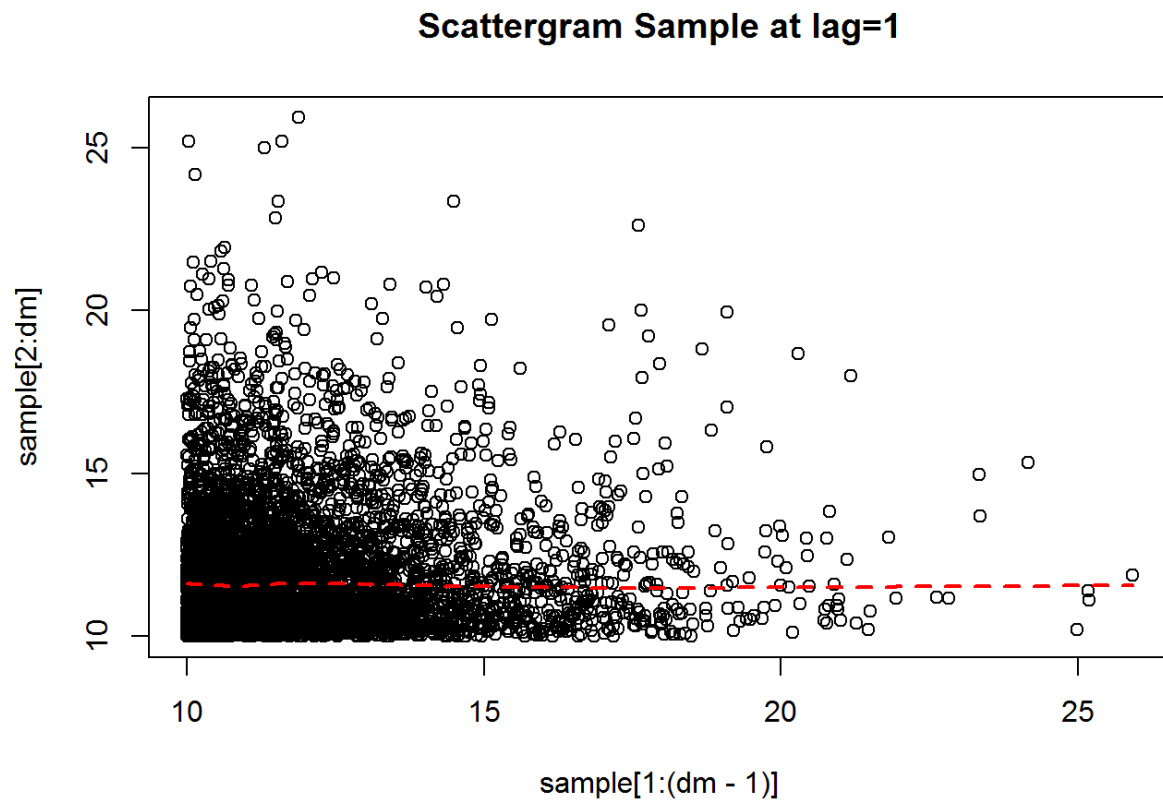
```
## [1] 1.0000000e+00 7.023715e-16 4.114827e+00 1.581992e+01 1.324174e+02
```

```
dm=5000 #sample size
```


1.2 Autocorrelation Function

- An scattergram for `data(1:(m-1))` vs `data(2:m)` is also valid and check for a flat smoother
- Default `scatterplot()` in `library(car)` contains linear adjustment and smoothers directly
- `acf()` is easy to interpret

```
library(car)
plot(sample[1:(dm-1)],sample[2:dm],main="Scattergram Sample at lag=1")
lines(lowess(sample[1:(dm-1)],sample[2:dm],f=0.5),col=2,lwd=2,lty=2) # Overlap smoother
```



```
scatterplot(sample[1:(dm-1)],sample[2:dm],main="Scattergram Sample at lag=1")
```




```
library(lmtest)
```

```
## Loading required package: zoo
##
## Attaching package: 'zoo'
##
## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric
```

```
dwtest(sample~1) # OK for residuals in library car and lmtest
```

```
##
## Durbin-Watson test
##
## data: sample ~ 1
## DW = 2.0122, p-value = 0.6674
## alternative hypothesis: true autocorrelation is greater than 0
```

```
## Also ?durbinWatsonTest # in library car
acf(sample)
```



2 Plot data

2.1 Histogram: Equal length intervals


Non suitable for distribution fitting Chi-squared Test

```
hist(sample,freq=F,breaks=100,main="Take a look to data...")  
m=mean(sample);std=sd(sample);m;std
```

```
## [1] 12.03262
```

```
## [1] 2.028707
```

```
lines(density(sample),col="blue")
```

Take a look to data...

Overlap some candidate distributions to fit data: normal (unlikely) and exponential (defined by rate parameter)

The exponential distribution with rate λ and location c has density $f(x) = \lambda * \exp(-\lambda(x - c))$ for $x > c$. The exponential cumulative distribution function with rate λ and location c is $F(x) = 1 - \exp(-\lambda(x-c))$ on $x > c$.

Theoretical moments for exponential distributions are:

1. $\text{Mean}(X) = E(X) = c + 1/\lambda$
2. $\text{Var}(X) = (1/\lambda)^2$

Location parameter c has to be estimated externally: for example, using the minimum, and for overlapped distributions should consider non-shifted distribution candidates.


```
#ussample<-sample-min(sample)+epsilon
ussample<-sample-10
hist(ussample,freq=F,breaks=100,main="Take a look to data...")
m=mean(ussample);std=sd(ussample);m;std
```

```
## [1] 2.032619
```

```
## [1] 2.028707
```

```
curve(dnorm(x,m,std),col="red",lwd=2,add=T)
curve(dexp(x,rate=1/m),col="green",lwd=2,add=T)
```

Take a look to data...



```
all.moments(ussample,4,central=F)
```

```
## [1] 1.000000 2.032619 8.246368 49.309399 380.113990
```

```
all.moments(ussample,4,central=T)
```

```
## [1] 1.000000e+00 -1.920242e-16 4.114827e+00 1.581992e+01 1.324174e+02
```

3 List of Candidate distributions

Formulate the list of candidate distributions: for distributions with shape parameter, plot the distribution for several shape parameters, using massive R plot, as the ones suggested in the following example, that takes a gamma distribution as possible candidate,

3.1 Plot pdf's for variable location = 0, scala = 1 for several shape parameter values

```
par( mflow = c( 3,2 ) )  
m.shape1 <-  rgamma(5000, shape=1, scale = 1)  
mean( m.shape1);sd( m.shape1);var( m.shape1)
```

```
## [1] 0.9884699
```

```
## [1] 0.9980591
```

```
## [1] 0.996122
```

```
hist( m.shape1, freq=F )  
x<-seq(0,max(m.shape1),100)  
curve(dgamma(x,shape=1,scale=1),col=2,add=T)  
  
m.shape0.5 <-  rgamma(5000, shape=0.5, scale = 1)  
mean( m.shape0.5);var( m.shape0.5)^0.5;var( m.shape0.5)
```

```
## [1] 0.4800706
```

```
## [1] 0.6602639
```

```
## [1] 0.4359484
```

```
hist( m.shape0.5, freq=F )  
x<-seq(0,max(m.shape0.5),100)  
curve(dgamma(x,shape=0.5,scale=1),col=2,add=T)  
  
m.shape2<-  rgamma(5000, shape=2, scale = 1)  
mean( m.shape2);var( m.shape2)^0.5;var( m.shape2)
```

```
## [1] 1.965469
```

```
## [1] 1.385176
```

```
## [1] 1.918712
```



```
hist( m.shape2, freq=F )
x<-seq(0,max(m.shape2),100)
curve(dgamma(x,shape=2,scale=1),col=2,add=T)

m.shape3<-  rgamma(5000, shape=3, scale = 1)
mean( m.shape3);var( m.shape3)^0.5;var( m.shape3)
```

```
## [1] 2.994269
```

```
## [1] 1.716285
```

```
## [1] 2.945635
```

```
hist( m.shape3, freq=F )
x<-seq(0,max(m.shape3),100)
curve(dgamma(x,shape=3,scale=1),col=2,add=T)

m.shape5<-  rgamma(5000, shape=5, scale = 1)
mean( m.shape5);var( m.shape5)^0.5;var( m.shape5)
```

```
## [1] 4.976568
```

```
## [1] 2.251865
```

```
## [1] 5.070896
```

```
hist( m.shape5, freq=F )
x<-seq(0,max(m.shape5),100)
curve(dgamma(x,shape=5,scale=1),col=2,add=T)

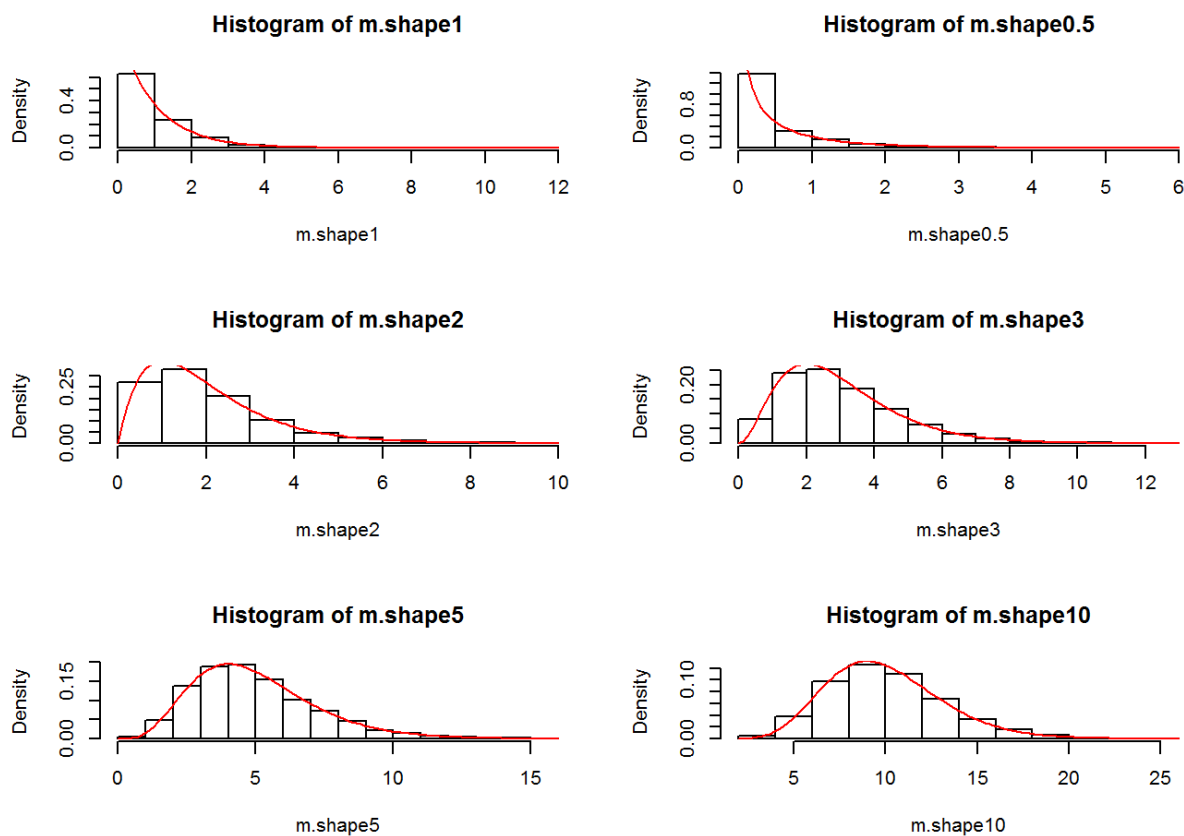
m.shape10 <-  rgamma(5000, shape=10, scale = 1)
mean( m.shape10);var( m.shape10)^0.5;var( m.shape10)
```

```
## [1] 10.04493
```

```
## [1] 3.192368
```

```
## [1] 10.19121
```

```
hist( m.shape10, freq=F )
x<-seq(0,max(m.shape10),100)
curve(dgamma(x,shape=10,scale=1),col=2,add=T)
```




3.2 Method descdist() in fitdistrplus package

```
library(fitdistrplus)
```

```
## Loading required package: MASS
```

```
descdist(ussample)
```

Cullen and Frey graph



```
## summary statistics
## -----
## min:  0.0003954973    max:  15.915
## median:  1.407226
## mean:  2.032619
## estimated sd:  2.028707
## estimated skewness:  1.895864
## estimated kurtosis:  7.826658
```

4 Fit distribution

To fit: use `fitdistr()` method in MASS package. Pay attention to supported distributions and how to refer to them (the name given by the method) and parameter names and meaning.

For discrete data use `goodfit()` method in vcd package: estimates and goodness of fit provided together

```
library(MASS)

res<-fitdistr(ussample,"weibull",lower=0.001)
res$estimate
```

```
##      shape      scale
## 0.9973711 2.0303540
```

```
# ?fitdistr    Check it
```

```
library(vcd) # For discrete data
```

```
## Loading required package: grid
```

```
# gf <- goodfit(d27x, type = "binomial", par = list(prob = 0.8, size = 7))
# summary(gf)
# plot(gf)

#Method fitdist() in fitdistrplus package
library(fitdistrplus)
fitg<-fitdist(ussample,"gamma")
fitw<-fitdist(ussample,"weibull")
```

5 Goodness of Fit

5.1 Histogram: Variable length intervals

- Histogram with breaks defined using quartiles of theoretical candidate distributions.
- Non Equal length intervals defined by empirical quartiles are more suitable for distribution fitting Chi-squared Test, since degrees of freedoms for Chi-squared Tests are guaranteed.

```
sequence<-seq(0,1,by=0.02)
qualist<-quantile(ussample,sequence)
sequence;qualist
```

```
## [1] 0.00 0.02 0.04 0.06 0.08 0.10 0.12 0.14 0.16 0.18 0.20 0.22 0.24
0.26
## [15] 0.28 0.30 0.32 0.34 0.36 0.38 0.40 0.42 0.44 0.46 0.48 0.50 0.52
0.54
## [29] 0.56 0.58 0.60 0.62 0.64 0.66 0.68 0.70 0.72 0.74 0.76 0.78 0.80
0.82
## [43] 0.84 0.86 0.88 0.90 0.92 0.94 0.96 0.98 1.00
```

```
##           0%           2%           4%           6%           8%
## 3.954973e-04 4.394407e-02 8.151098e-02 1.250477e-01 1.641681e-01
##           10%          12%          14%          16%          18%
## 2.060015e-01 2.529657e-01 2.968658e-01 3.513416e-01 3.987952e-01
##           20%          22%          24%          26%          28%
## 4.536233e-01 5.102732e-01 5.624885e-01 6.199457e-01 6.694348e-01
##           30%          32%          34%          36%          38%
## 7.195240e-01 7.775978e-01 8.409661e-01 9.041210e-01 9.738403e-01
##           40%          42%          44%          46%          48%
## 1.038386e+00 1.104926e+00 1.180532e+00 1.254849e+00 1.322035e+00
##           50%          52%          54%          56%          58%
## 1.407226e+00 1.492833e+00 1.571155e+00 1.654192e+00 1.747934e+00
##           60%          62%          64%          66%          68%
## 1.855825e+00 1.964084e+00 2.073619e+00 2.191610e+00 2.321625e+00
##           70%          72%          74%          76%          78%
## 2.440747e+00 2.585901e+00 2.749210e+00 2.899563e+00 3.070045e+00
##           80%          82%          84%          86%          88%
## 3.250987e+00 3.438255e+00 3.687353e+00 3.961713e+00 4.300690e+00
##           90%          92%          94%          96%          98%
## 4.691573e+00 5.162643e+00 5.862118e+00 6.690737e+00 7.969253e+00
##           100%
## 1.591500e+01
```


```
hist(ussample,freq=F,breaks=qualist,main="Take a look to data...")
m=mean(ussample);std=sd(ussample);m;std
```

```
## [1] 2.032619
```

```
## [1] 2.028707
```

```
curve(dnorm(x,m,std),col="red",lwd=2,add=T)
curve(dexp(x,rate=1/m),col="green",lwd=2,add=T)
lines(density(sample),col="blue")
lines(density(ussample),col="cyan")
```

Take a look to data...



Method fitdist() in fitdistrplus package

```
library(fitdistrplus)
# descdist(ussample)
fitg<-fitdist(ussample,"gamma")
fitw<-fitdist(ussample,"weibull")

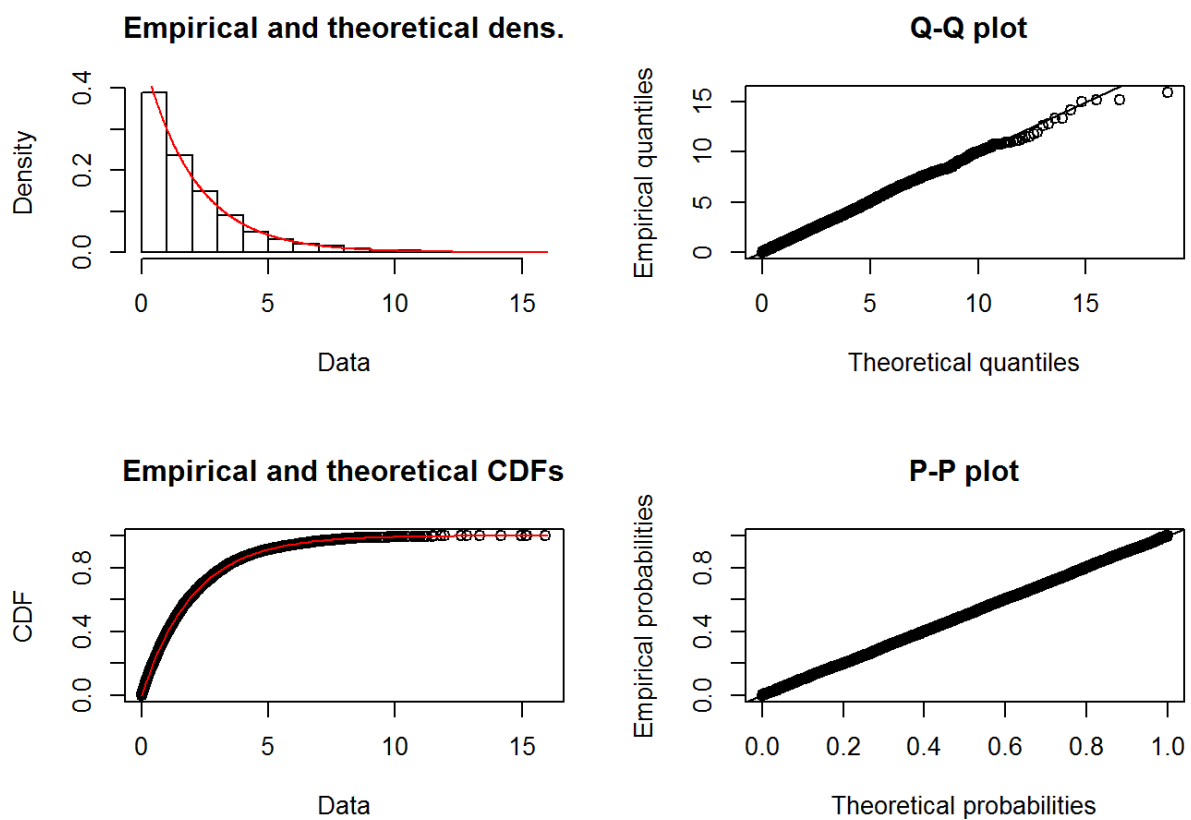
summary(fitg)
```

```
## Fitting of the distribution ' gamma ' by maximum likelihood
## Parameters :
##      estimate Std. Error
## shape 0.9950789 0.01751549
## rate  0.4895212 0.01106380
## Loglikelihood: -8546.587   AIC:  17097.17   BIC:  17110.21
## Correlation matrix:
##      shape      rate
## shape 1.0000000 0.7788053
## rate  0.7788053 1.0000000
```


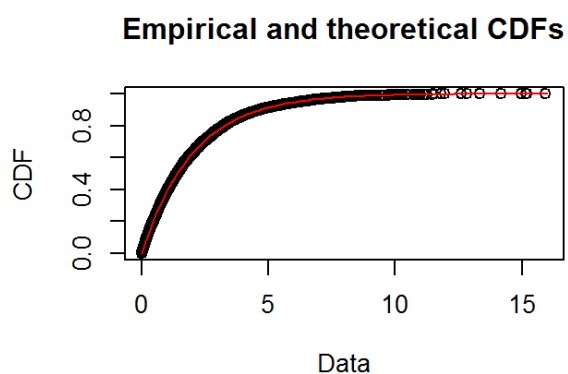


```
summary(fitw)
```

```
## Fitting of the distribution ' weibull ' by maximum likelihood
## Parameters :
##      estimate Std. Error
## shape 0.9973664 0.01101032
## scale 2.0301398 0.03030799
## Loglikelihood: -8546.598   AIC: 17097.2   BIC: 17110.23
## Correlation matrix:
##           shape      scale
## shape 1.0000000 0.3131775
## scale 0.3131775 1.0000000
```


```
plot(fitg)
```



```
plot(fitg, demp = TRUE)
```




```
cdfcomp(list(fitg,fitw), legendtext=c("Gamma","Weibull"))
```





```
denscomp(list(fitg,fitw), legendtext=c("Gamma","Weibull"))
```


Histogram and theoretical densities



```
ppcomp(list(fitg,fitw), legendtext=c("Gamma","Weibull"))
```

P-P plot

```
qqcomp(list(fitg,fitw), legendtext=c("Gamma","Weibull"))
```

Q-Q plot

```
gofstat(list(fitg,fitw), fitnames=c("Gamma","Weibull"))
```

```
## Goodness-of-fit statistics
##
##           Gamma      Weibull
## Kolmogorov-Smirnov statistic 0.005389694 0.005259483
## Cramer-von Mises statistic  0.016380448 0.016222121
## Anderson-Darling statistic  0.176133193 0.175976428
##
## Goodness-of-fit criteria
##
##           Gamma      Weibull
## Aikake's Information Criterion 17097.17 17097.20
## Bayesian Information Criterion 17110.21 17110.23
```

5.2 Validate matching between empirical moments and theoretical moments


For each candidate distributions calculate up to degree 4 theoretical moments and check central and absolute empirical moments. Previously, you have to estimate parameters and calculate theoretical moments, using estimated parameters. Good matching should exists for any of the candidate distributions between theoretical and empirical moments.

5.3 Goodness of Fit: Graphical Assessment

5.3.1 Known shape

Location and scale parameter estimates are returned as coefficient of linear regression in QQPlot. Check versus fitdistr estimates for distribution parameters

```
xx <- qweibull( ppoints( sample ) , shape=1)
yy<- sort( sample )
mm <- lm( yy ~ xx )
plot(xx,yy,main="QQPlot")
abline(mm,col="red", lwd=2, lty=2)
```



```
cors<-cor(yy,xx);cors
```


```
## [1] 0.9994532
```

5.3.2 UNKnown shape

Use standardized distributions - Identifies shape giving the best fit (alternative to ML estimation). As a subproduct location and scale parameters are also estimated, so you do not need to unshift your data.

```
rang <- seq(0.5,2,length=50)
cors<-rep(0,50)

for (i in 1:50)
{
xx <- qweibull( ppoints( sample ) , shape=rang[ i ] )
yy<- sort( sample )
mm <- lm( yy ~ xx )
cors[ i ]<-cor(yy,xx);cors
}
plot(rang,cors)
```



```
# Goodness of Fit: Histogram needs estimates for parameters
sequence<-seq(0,1,by=0.02)
qualist<-qweibull(shape=res$estimate[1],scale=res$estimate[2],sequence)
sequence;qualist
```

```
## [1] 0.00 0.02 0.04 0.06 0.08 0.10 0.12 0.14 0.16 0.18 0.20 0.22 0.24
0.26
## [15] 0.28 0.30 0.32 0.34 0.36 0.38 0.40 0.42 0.44 0.46 0.48 0.50 0.52
0.54
## [29] 0.56 0.58 0.60 0.62 0.64 0.66 0.68 0.70 0.72 0.74 0.76 0.78 0.80
0.82
## [43] 0.84 0.86 0.88 0.90 0.92 0.94 0.96 0.98 1.00
```

```
## [1] 0.00000000 0.04059894 0.08218727 0.12471091 0.16818922 0.21265402
## [7] 0.25814354 0.30470080 0.35237305 0.40121176 0.45127272 0.50261636
## [13] 0.55530810 0.60941881 0.66502535 0.72221122 0.78106732 0.84169277
## [19] 0.90419592 0.96869554 1.03532213 1.10421952 1.17554676 1.24948028
## [25] 1.32621656 1.40597523 1.48900288 1.57557761 1.66601471 1.76067348
## [31] 1.85996592 1.96436748 2.07443082 2.19080353 2.31425107 2.44568735
## [37] 2.58621568 2.73718507 2.90026900 3.07757865 3.27183018 3.48660053
## [43] 3.72673446 3.99902496 4.31342174 4.68535172 5.14066239 5.72781537
## [49] 6.55562671 7.97140072          Inf
```

```
qualist[51]<-max(ussample)
```


```
hist(ussample,freq=F,breaks=qualist,main="Take a look to data...")  
m=mean(ussample);std=sd(ussample);m;std
```

```
## [1] 2.032619
```

```
## [1] 2.028707
```

```
curve(dnorm(x,m,std),col="red",lwd=2,add=T)  
curve(dexp(x,rate=1/m),col="green",lwd=2,add=T)  
lines(density(sample),col="blue")  
lines(density(ussample),col="cyan")
```

Take a look to data...



5.4 Goodness of Fit: Tests

5.4.1 Kolmogorov-Smirnoff Test for continuous data

Beware of using the proper names in R for distribution parameters

```
#ks.test(ussample,"pweibull") # Doesn't work
ks.test(ussample,"pweibull", scale=res$estimate[2], shape=res$estimate[1]
)
```

```
##
## One-sample Kolmogorov-Smirnov test
##
## data:  ussample
## D = 0.0052914, p-value = 0.999
## alternative hypothesis: two-sided
```

5.4.2 Cramer von Miess Test for discrete data

For discrete data (discrete version of KS Test). library(dgof) includes cvm.test() Cramer von Miess test, discrete version of KS Test.

5.4.3 Chi Squared Test

Chi Squared Test - It requires manual programming using non-constant length intervals (defined by quartiles). Valid for discrete or continuous data.

```
sequence<-seq(0,1,by=0.02)
qualist<-qweibull(shape=res$estimate[1],scale=res$estimate[2],sequence)
sequence;qualist
```

```
## [1] 0.00 0.02 0.04 0.06 0.08 0.10 0.12 0.14 0.16 0.18 0.20 0.22 0.24
0.26
## [15] 0.28 0.30 0.32 0.34 0.36 0.38 0.40 0.42 0.44 0.46 0.48 0.50 0.52
0.54
## [29] 0.56 0.58 0.60 0.62 0.64 0.66 0.68 0.70 0.72 0.74 0.76 0.78 0.80
0.82
## [43] 0.84 0.86 0.88 0.90 0.92 0.94 0.96 0.98 1.00
```

```
## [1] 0.00000000 0.04059894 0.08218727 0.12471091 0.16818922 0.21265402
## [7] 0.25814354 0.30470080 0.35237305 0.40121176 0.45127272 0.50261636
## [13] 0.55530810 0.60941881 0.66502535 0.72221122 0.78106732 0.84169277
## [19] 0.90419592 0.96869554 1.03532213 1.10421952 1.17554676 1.24948028
## [25] 1.32621656 1.40597523 1.48900288 1.57557761 1.66601471 1.76067348
## [31] 1.85996592 1.96436748 2.07443082 2.19080353 2.31425107 2.44568735
## [37] 2.58621568 2.73718507 2.90026900 3.07757865 3.27183018 3.48660053
## [43] 3.72673446 3.99902496 4.31342174 4.68535172 5.14066239 5.72781537
## [49] 6.55562671 7.97140072          Inf
```

```

qualist[51]<-max(ussample)
dsample<-cut(ussample,breaks=qualist)
iobs<-as.vector(table(dsample))  # Obs in the groups defined by nb intervals-percentiles
pexp<-as.vector(rep(1/50,50))    # Expected probability for each groups: sample size/nb intervals
iexp<-length(ussample)*pexp
iobs;pexp;iexp

```

```

## [1] 91 112 96 110 99 110 97 86 102 92 90 105 93 109 112 104 9
5
## [18] 97 91 107 101 96 94 117 93 97 110 115 91 94 95 99 99 9
3
## [35] 111 98 93 107 104 105 114 90 99 96 90 94 92 96 119 100

```

```

## [1] 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02
0.02
## [15] 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02
0.02
## [29] 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02
0.02
## [43] 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02 0.02

```

```

## [1] 100 100 100 100 100 100 100 100 100 100 100 100 100 100 100 100 100 10
0
## [18] 100 100 100 100 100 100 100 100 100 100 100 100 100 100 100 100 100 10
0
## [35] 100 100 100 100 100 100 100 100 100 100 100 100 100 100 100 100 100

```

```

X2<-sum(((iobs-iexp)^2)/iexp);X2

```

```

## [1] 33.66

```

```

#chisq.test( sample, "pweibull", shape= 2, scale=1 )  #Doesn't work
chisq.test( iobs,p=pexp ) # Works

```

```

##
## Chi-squared test for given probabilities
##
## data: iobs
## X-squared = 33.66, df = 49, p-value = 0.9535

```