

REPORT

Data Cleaning:

After downloading the data from UCI, I cleaned it to fit into the linear regression model.

At first I have selected the numerical type columns from the dataset:-

Selecting the numerical type columns

```
In [4]: data1=data.select_dtypes(include=np.number)
```

Then I count all NA values from the dataset.

Counting Null values per column

```
In [5]: data1.isna().sum()
```

```
Out[5]: Id                0
        MSSubClass         0
        LotFrontage       259
        LotArea            0
        OverallQual        0
        OverallCond        0
        YearBuilt          0
        YearRemodAdd        0
        MasVnrArea         8
        BsmtFinSF1         0
        BsmtFinSF2         0
        BsmtUnfSF          0
        TotalBsmtSF        0
        1stFlrSF           0
        2ndFlrSF           0
        LowQualFinSF       0
        GrLivArea          0
        BsmtFullBath        0
        BsmtHalfBath        0
        FullBath            0
        HalfBath            0
        BedroomAbvGr        0
        KitchenAbvGr        0
        TotRmsAbvGrd        0
        Fireplaces         0
        GarageYrBlt        81
        GarageCars          0
        GarageArea          0
        WoodDeckSF          0
        OpenPorchSF         0
        EnclosedPorch        0
        3SsnPorch           0
        ScreenPorch         0
        PoolArea            0
        MiscVal             0
        MoSold              0
        YrSold              0
        SalePrice           0
        dtype: int64
```

Since only few columns have NA values so I dropped them.

Removing NA Values


```
In [6]: df = data1.dropna()
        #d.isna().sum()#d.isna().sum()
```

Then I have plotted the heat map of the correlation matrix

Create correlation matrix and the heatmap

```
In [7]: corr = df.corr()
        plt.subplots(figsize=(15,10))
        ax = sns.heatmap(corr, cmap="BuPu", vmax=0.7, square=True)
        ax.set_title("Correlations", fontsize = 18)
```

```
Out[7]: Text(0.5, 1.0, 'Correlations')
```



Using the correlation matrix heat map, I have selected the most suitable columns for linear regression.

Then I have converted the data frame into two numpy arrays, named them X and Y.

Converting the dataframe into two numpy arrays

```
: X = np.array(df[['GarageCars',  
                  'GarageArea',  
                  'GrLivArea',  
                  'OverallQual',  
                  'TotRmsAbvGrd',  
                  '1stFlrSF',  
                  'FullBath',  
                  'WoodDeckSF',  
                  'OpenPorchSF']])  
  
y = np.array(df['SalePrice'])
```

OLS Regression Results

Ordinary Least Squares

In [34]: `%matplotlib inline`

In [35]: `res = sm.OLS(y_train, X_train).fit()
print(res.summary())`

```

OLS Regression Results
=====
Dep. Variable:          y      R-squared (uncentered):      0.955
Model:                  OLS    Adj. R-squared (uncentered):    0.954
Method:                 Least Squares    F-statistic:          1808.
Date:                  Thu, 17 Mar 2022    Prob (F-statistic):    0.00
Time:                  15:12:12    Log-Likelihood:      -9479.8
No. Observations:      784    AIC:                  1.898e+04
Df Residuals:          775    BIC:                  1.902e+04
Df Model:              9
Covariance Type:       nonrobust
=====
                    coef    std err          t      P>|t|      [0.025      0.975]
-----
x1             1.31e+04    4781.282        2.740      0.006     3717.256     2.25e+04
x2              37.3806      15.702        2.381      0.018         6.557         68.205
x3              84.4390       5.988       14.102      0.000         72.685         96.193
x4             1.553e+04    1392.134       11.157      0.000        1.28e+04        1.83e+04
x5             -1.517e+04    1454.407      -10.430      0.000       -1.8e+04       -1.23e+04
x6              19.7530       5.271        3.747      0.000         9.405         30.101
x7             -8185.1417    3920.950       -2.088      0.037       -1.59e+04       -488.201
x8              58.4710      13.432        4.353      0.000         32.104         84.838
x9              67.0294      25.666        2.612      0.009         16.646        117.413
=====
Omnibus:                 308.490    Durbin-Watson:           2.068
Prob(Omnibus):            0.000    Jarque-Bera (JB):        7148.142
Skew:                     1.222    Prob(JB):                 0.00
Kurtosis:                 17.589    Cond. No.                 6.58e+03
=====

```


Notes:

- [1] R^2 is computed without centering (uncentered) since the model does not contain a constant.
- [2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Batched Gradient Descent

Train Score: 0.7190043786890068


Test Score: 0.6269585571709697



Sequential Gradient Descent

Train Score: 0.7193427721554584

Test Score: 0.6273421816020116



In this case Sequential Gradient Descent takes more iterations and give slightly better result than Batched Gradient Descent.